

PRIYA SHARMA

Senior ML Engineer | priya.sharma@example.com | +91 98765 43210
Bangalore, India | linkedin.com/in/priyasharma | github.com/priyasharma
Current CTC: INR 32 LPA | Notice Period: 60 days

PROFESSIONAL SUMMARY

Senior Machine Learning Engineer with 6 years of production experience shipping LLM-based systems, distributed training pipelines, and recommendation systems. Strong track record of taking models from research to production at scale (100+ RPS, 99.9% uptime).

SKILLS

Python, PyTorch, TensorFlow, JAX, Hugging Face Transformers, LangChain, AWS SageMaker, GCP Vertex AI, Docker, Kubernetes, Ray, MLflow, DVC, Triton Inference Server, FastAPI, distributed training, RAG, vector DBs (Pinecone, Weaviate), LLM fine-tuning, MLOps, CI/CD

EXPERIENCE

Senior ML Engineer | Acme AI | 2022 - Present

- Built production RAG system over 8M documents; latency p95 under 200ms at 200 RPS
- Led distributed training migration to 32-GPU clusters; cut training time 60%
- Owned cost optimization: migrated inference to Triton, reduced spend 40%
- Fine-tuned 7B-parameter LLMs for internal Q&A; precision 92%

ML Engineer | DataBay Technologies | 2019 - 2022

- Built feature store and model registry; supports 50+ production models
- Shipped recommendation system to 12M users; lift +14% on engagement
- Designed A/B testing framework for ML models

EDUCATION

B.Tech Computer Science | IIT Bombay | 2019 | CGPA 9.2/10

CERTIFICATIONS

AWS Certified Machine Learning - Specialty (2023)
DeepLearning.AI Deep Learning Specialization (2020)

PUBLICATIONS

"Efficient Retrieval Augmentation for Multi-Tenant LLMs" - ICML Workshop, 2024

PROJECTS

OpenRAG: open-source RAG framework (1.2k GitHub stars)
ML Bias Audit Toolkit: detected gender/caste bias signals in 4 production models